

AI & RAG EXPLAINED

What is Artificial Intelligence?

Artificial Intelligence (AI) refers to the ability of machines to perform tasks that typically require human intelligence. These tasks include learning, reasoning, problem-solving, and understanding language.

AI systems are widely used in applications like chatbots, recommendation systems, and self-driving cars.

What is Python?

Python is a high-level, easy-to-learn programming language widely used in web development, data science, and artificial intelligence.

Its simple syntax and powerful libraries make it one of the most popular languages for building AI applications.

What is RAG (Retrieval-Augmented Generation)?

RAG stands for Retrieval-Augmented Generation.

It is a technique that combines:

1. Retrieval (searching relevant data)
2. Generation (creating answers using AI)

Instead of relying only on pre-trained knowledge, RAG allows AI systems to fetch relevant information from external sources like documents, PDFs, or databases.

This makes the responses more accurate and up-to-date.

How RAG Works

The RAG process involves the following steps:

1. A user asks a question.
2. The system searches a database for relevant information.
3. The retrieved data is sent to a language model.
4. The language model generates a final answer based on that data.

This ensures that answers are grounded in actual data instead of guesses.

What are Embeddings?

Embeddings are numerical representations of text.

They convert words and sentences into vectors (lists of numbers) so that machines can understand meaning.

For example:

- "car" and "vehicle" will have similar embeddings.
- "car" and "banana" will have very different embeddings.

What is FAISS?

FAISS (Facebook AI Similarity Search) is a library used to search similar data quickly.

It stores embeddings and helps retrieve the most relevant information based on similarity.

FAISS is commonly used in RAG systems to find the best matching text chunks.

What is a Vector Database?

A vector database stores data in the form of embeddings.

Instead of searching by exact words, it searches by meaning.

This allows AI systems to find relevant information even if the wording is different.

What is an LLM?

A Large Language Model (LLM) is an AI model trained on massive amounts of text data.

It can generate human-like responses, answer questions, summarize text, and more.

Examples include ChatGPT and GPT-based models.

Why Use RAG?

RAG solves major limitations of AI:

- It provides up-to-date information
- It reduces hallucinations (wrong answers)

- It allows AI to use private or custom data

This makes it ideal for building real-world AI applications.

Real-World Use Cases of RAG

RAG is used in many applications, such as:

- Customer support chatbots
 - Document-based Q&A systems
 - Legal and medical assistants
 - Internal company knowledge systems
-

Limitations of RAG

While RAG is powerful, it has some limitations:

- Requires good quality data
 - Retrieval may sometimes miss relevant information
 - Performance depends on chunking and embeddings
-

Future of AI with RAG

RAG is becoming a standard approach for building AI applications that need accurate and reliable information.

It allows developers to create AI systems that are not just intelligent, but also trustworthy.